Telstra Data Normalisation

Team Koala

Luke
Joel
Ishan
Nikunj
Daniel

The Problem: Data Normalisation

Thousands of devices

Variety of formats

Security analysis requires data normalization

Must Normalize to standard schema

The Process: Data Normalisation

1

Constructing Mappings

2

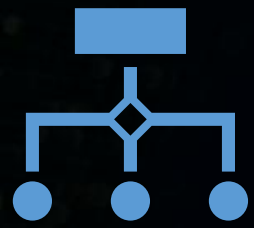

Normalising Live Data

1. Current Constructing Mappings

Online Documentation

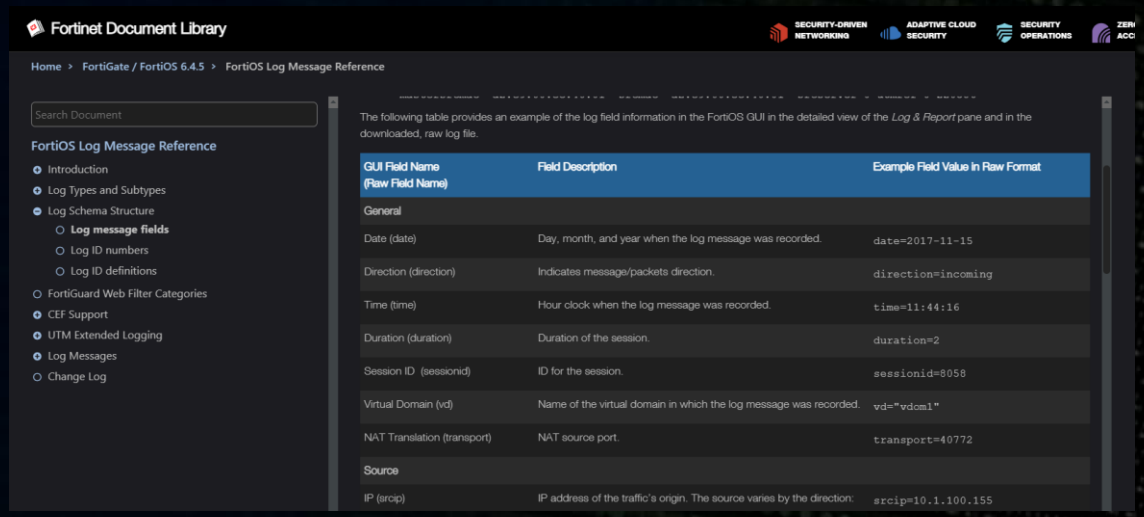

Sample Log

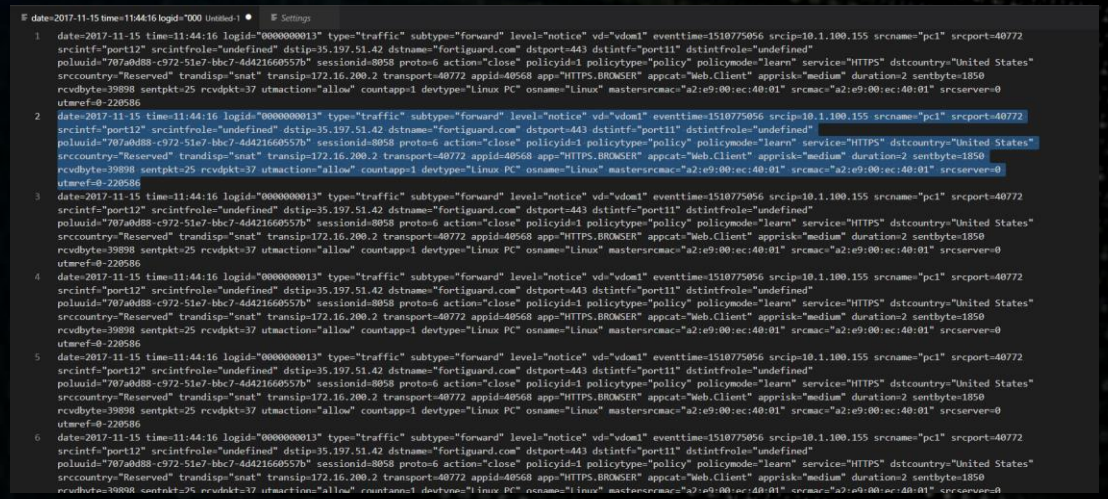

Internal Documentation

ECS Field	Syslog/Dark Trace Field Name	Example Output	Comments
@timestamp	"@timestamp"		This should be the time the event was parsed into Elastic, in UTC
tenant.name	N/A		Manually set in the router, don't change after this
event.start	syslog header timestamp	<165>**Nov 22 09:02:32**10.107.250.1 {"creationTime":...	This should be the original device/event timestamp from the syslog header, in UTC
event.created	creationTime	1542877373000 (*requires appropriate conversion*)	This should be the time of the event from within the contents of the event, in UTC The timestamp that the record of the breach was created. This is distinct from the "time" field.
event.timezone	N/A	*specific to customer*	This should be the original device/event timezone eg "AEST" and may need to be manually specified at parsing time based on customer

2. Normalising Live Data

Internal Documentation

ECS Field	Syslog/Dark Trace Field Name	Ex	
@timestamp	"@timestan "		
tenant.name			
event.start	sys he tin		
event.created	creation t	*(requ appropria conversion)*	stamp record creat the "	each was distinct from
event.timezone	N/A	*specific to customer*	This d be the original device/event timezone eg "AEST" and may need to be manually specified at parsing time based on customer

Automation

pping Configuration

```
acking used along with src_ip and dst_ip fields

{srcip}" }

tion => "append"

uns" ]

ove unnecessary fields to keep ES memory cache from filli
e you would want to comment certain types or tags out if trying

e_field => [ "host", "received_at", "received_from", "syslog_hostname", "s
_tag => "syslog"
        }
    }
}
# Send output to local elasticsearch instance
# Change to one of the other modes and comment out below if needed
output {
    elasticsearch_http {
        host => "127.0.0.1"
        flush_size => 1
        template_overwrite => true
        manage_template => true
        template => "/opt/logstash/lib/logstash/outputs/elasticsearch/elasticsearch-template.js
    }
}
```

Motivation

Simplify the process

Reduce time and effort

Improve accuracy

Easier Automation

The Process: Data Normalisation

1

Our Solution

Constructing Mappings

2

Extension

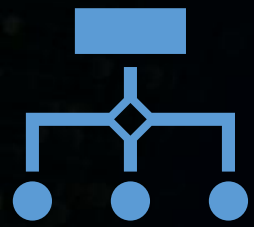

Normalising Live Data

Our Solution: Core Functionality

Normalising input and output fields

source_ip ➡ source_destination

UINT32

Typing of Fields

dateTime as EPOCH
dateTime as UTC

(.*)

Types represent regular expressions

IPv4
\b((25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)(\.|$)){4}\b

149.167.143.3: Match
Bananas: No Match

Our Solution: UX Functionality

Suggestion of mappings

Cloning existing mappings and formats

Validation against sample data

Extensions

Generating Splunk/Elastic config files automatically

Read and detect different log formats

Integrated Ingestion of Log files

API access of documents

Generating validation code for continually integration

Demo

1. Creating a Type
2. Creating an Input Format
3. Validating Input Format
4. Creating a Normalisation Format

Questions